AI Adoption & Agent Development for ARG Companies

A practical path to faster turnarounds, freed senior capacity, and governed AI — built for asset-based lending and restructuring work

Prepared by ITECS for ARG Companies leadership · June 2026

Scope: (1) AI adoption & workforce training · (2) Custom AI agent development

Audience: ARG principals and practice leads (Partners, Recovery, Industrial, Asset Recovery, Transaction Services, AR Management)

Executive summary

ARG's product is senior judgment delivered fast: field examinations, comprehensive and 13-week cash flow models, NOLV and M&E appraisals, wind-down reporting, AR analysis, and the monthly commodity report. Every one of those deliverables is built from the same raw materials — borrower financials, AR/AP agings, inventory and equipment listings, general ledgers, bank statements, trial balances, and borrowing base certificates — and today a seasoned professional spends a large share of each engagement *assembling and normalizing* that data before the judgment work even begins.

That assembly layer is exactly where governed AI creates leverage. The opportunity for ARG is twofold:

- **AI adoption & training** — get every professional fluent and safe with the leading AI tools (Claude, ChatGPT, and Microsoft Copilot), with a clear acceptable-use policy that protects confidential lender and borrower data.
- **Custom AI agents** — purpose-built agents that do the repetitive collateral analysis, roll-forwards, reconciliations, and first-draft reporting, so principals review and decide instead of build.

The payoff ARG should expect: **shorter exam and model turnaround for lenders, more engagements handled by the same senior bench, and a defensible, audit-ready trail behind every number.** ARG's own model already proves the thesis — the firm notes that prior reports supply *roughly 80% of what's needed* to build a comprehensive cash flow model. Agents are how that 80% gets assembled in hours instead of days.

Every agent in this document follows three principles:

- **Tied to a real ARG deliverable** — an exam, a model, an appraisal, a report. No “AI for its own sake.”
- **Read broadly, act carefully** — agents read across borrower data and *draft and recommend*; they never send communications, alter collateral records, or finalize a valuation without a professional's review and an audit trail.
- **Confidential by design** — built for a firm that handles distressed-company financials, lender collateral data, and bankruptcy-sensitive information.

Part 1 — AI Adoption & Training

The fastest, lowest-risk value for a professional-services firm comes from getting the whole team productive and safe with AI tools they can use today — inside Excel models, Word reports, Outlook, and the desktop apps.

1.1 The three platforms, side by side

ARG's staff will get the most from the three leading platforms, each with a desktop app plus business/enterprise tiers. The right mix depends on ARG's Microsoft 365 footprint and where the confidential data lives.

Plan / Tier	Anthropic — Claude	OpenAI — ChatGPT	Microsoft 365 Copilot
Free / included	Claude Free — $0	ChatGPT Free — $0 [1]	Copilot Chat — included [2]
Individual	Pro — $20/mo [3]	Go — $8 · Plus — $20/mo	— (organization-licensed)
Power user	Max — $100–$200/mo	Pro — $200/mo	—
Business (per seat)	Team — $20–$25/seat [4]	Business — $20–$25/seat [4]	Copilot add-on — $21–$30/user [5]
Enterprise	Custom	Custom	Custom (+ usage [6])

[1] ChatGPT Free is ad-supported in the US (since Feb 2026). [2] Copilot Chat is included at no added cost with eligible M365 plans; secure chat with enterprise data protection. [3] Claude Pro is ~$17/mo billed annually. [4] Per-seat business pricing: lower figure billed annually, higher billed monthly; ChatGPT Business moved to $20 (annual)/$25 (monthly) in April 2026, and Claude Team also offers premium seats at $100–$125/seat. [5] Microsoft 365 Copilot is an add-on to a qualifying M365 plan; rate varies by agreement (Business ~$21, Enterprise $30) — the all-in cost including the required base M365 license typically runs ~$34–$87/user/mo. [6] Copilot Studio agents are billed on consumption (capacity/credits), not per seat. List pricing in USD, current as of June 27, 2026 and subject to change. ITECS confirms exact pricing against ARG's M365 agreement.

How they map to ARG's work:

- **Microsoft 365 Copilot** — strongest where ARG already lives: Excel (cash flow models, roll-forwards, borrowing base math), Word (Executive Reports, workpapers), Outlook, and Teams, grounded in ARG's own documents. Best fit if ARG is on a qualifying M365 plan.
- **Claude (desktop app)** — excellent for long-document analysis, financial-statement and contract reading, structured reasoning over messy borrower files, and drafting reports in ARG's voice. Strong analytical writing and large-context handling.
- **ChatGPT (desktop app)** — versatile general assistant for research, summarization, ad-hoc analysis, and quick drafting; broad ecosystem and custom GPTs.

ITECS is platform-agnostic — we recommend the right fit per use case rather than locking ARG into one vendor.

1.2 The security questions that matter for a financial-advisory firm

ARG handles lender collateral data, distressed-borrower financials, and bankruptcy-sensitive material. The three questions every principal will ask:

1. **Is our data used to train the AI or exposed to outsiders?** No — on business/enterprise accounts, prompts, the files the AI reads, and responses stay inside the provider's protected boundary, are not used to train the models, and are not shared with other customers. Microsoft 365 Copilot keeps this within ARG's M365 service boundary.

2. **Can the AI see data a staffer shouldn't?** It only surfaces what that user already has permission to open — it enforces existing permissions rather than bypassing them. The real risk is internal oversharing (e.g., an over-broad SharePoint site), which ITECS reviews and tightens before rollout.
3. **How is a business account different from free/consumer AI?** Consumer tools may use inputs to improve their models and sit outside ARG's control. Business accounts contractually isolate ARG's data, add audit logging and admin governance, and keep content in a protected boundary. Rule of thumb: confidential engagement data goes only into approved business tools.

The single biggest near-term risk is **shadow AI** — staff pasting borrower financials into free consumer tools. ITECS addresses this with a discovery scan, a clear acceptable-use policy, and approved business-tier tooling.

1.3 The adoption path (phased, decision points at each step)

Phase	What it covers	Key deliverables	List price
1. Executive AI Orientation	Half-day leadership session: how AI works, where it creates value in ARG's specific workflows, governance overview, Q&A. Tailored to ARG.	Landscape briefing across Claude/ChatGPT/Copilot; ARG-relevant use-case discussion	$3,500 / session
2. AI Readiness Assessment	1–2 week engagement: current-state audit, shadow-AI exposure scan, prioritized use-case map, recommended platform stack, 12-month roadmap, executive readout (incl. 90-min leadership workshop).	Written roadmap; use-case prioritization; platform recommendation; governance starter	$6,500
Foundation: Acceptable Use Policy	Custom AI Acceptable Use Policy + governance framework + employee-handbook insert + training deck. Coordinated with ARG's counsel where applicable.	AUP package covering all platforms in scope	$4,500
Foundation: Shadow AI Discovery	Endpoint + M365 scan for unsanctioned AI usage, risk-ranked findings, remediation playbook, board-ready summary.	Shadow AI risk report	$5,500
3. AI Pilot Implementation — Small	Primary platform tenant/workspace setup, 5-template prompt library, 1 training session, 30-day optimization. 3–5 users, one use case.	Production-ready config; prompt library; training	$12,500
3. AI Pilot Implementation — Production	Multi-use-case expansion, role-based prompt libraries, audit-trail documentation, 2 training sessions, 60-day optimization. 10–25 users.	Everything in Small, scaled	$21,500
Enablement: Internal AI Champion	4–6 week structured program training one ARG employee to be the firm's internal AI lead: curriculum + weekly coaching + project shadowing + handoff docs.	Internal champion + documentation	$8,500–$12,000

Phase	What it covers	Key deliverables	List price
4. Managed AI Services	Ongoing optimization, prompt-library maintenance, platform monitoring, QBRs, standing access to ITECS's AI team.	Monthly advisory + maintenance	$1,950 / $2,650 / $3,500 per mo [7]

[7] *Managed AI tiers: Starter (≤10 users, 4 hrs), Standard (11–25 users, 6 hrs), Plus (26–50 users, 10 hrs). These are adoption/advisory retainers; production agents that need monitoring and tuning are covered by Managed AI — Agent Operations (Part 2).*

Recommended sequence for ARG: Orientation → Readiness Assessment → AUP + Shadow AI (run together) → a Production pilot in the highest-volume practice (Transaction Services is the natural first home) → Internal Champion + Managed AI to sustain it.

1.4 Training delivery & hourly work

Hands-on enablement, custom prompt libraries, and bespoke configuration outside the productized pilots are billed by tier:

Tier	Role	What it covers	Rate
Tier 1	AI Implementer	Tenant/workspace setup, prompt-library build, employee training delivery, documentation, champion coaching	$245/hr
Tier 2	AI Specialist	Custom agent/connector/MCP development, integration, prompt engineering, policy authoring, workflow design	$295/hr
Tier 3	AI Strategist	Executive workshops, AI roadmapping, fractional Head of AI, AI security architecture, governance & audit support	$375/hr

Part 2 — Custom AI Agent Development

This is the higher-leverage half. Below is the **full menu** of agent opportunities, grouped by ARG practice area, followed by a flagship deep-dive, pricing, governance, and an ROI framing. Each agent lists what it does in ARG's context, the metric it moves, an illustrative value, and the data it draws on.

How agent builds are priced. ARG's agents range from light, bounded assistants to multi-source financial-workpaper systems, so ITECS scopes them in phases rather than at one flat fee. A short, fixed-fee discovery step defines the data, architecture, human-review workflow, and a not-to-exceed for the build — and is credited toward the build if ARG proceeds. The readiness assessment ranks the menu by payback so the first build has the cleanest data and clearest return.

Build stage	ITECS price	When it applies
Agent Discovery & Technical Specification	$4,500–$7,500	Data/source review, architecture, risk & review model, fixed build estimate. Credited toward the build.

Build stage	ITECS price	When it applies
Proof of Concept / Prototype	\$8,000–\$18,000	One data source or workflow, limited users, no production SLA — proves feasibility.
Single-Workflow Production Agent	\$18,000–\$35,000	One primary workflow, deployed and documented, human-in-the-loop.
Integrated / Financial Workpaper Agent	\$35,000–\$75,000	Multi-source ingestion and reconciliation, audit trail, validation, exception handling, reviewer workflow.
Multi-Agent System / Process Redesign	\$55,000–\$120,000	Orchestrated multi-agent or end-to-end workflow redesign with integrations and monitoring.

Where ARG's agents typically land: the bounded ones — Commodity Report (E1), Executive Briefing (E3), and M&E Comparable-Research (C2) — sit in the single-workflow band; most collateral and modeling agents — BBC (A2), Portfolio Monitoring (A3), 13-Week Cash Flow (B2), NOLV Accelerator (C1), Wind-Down Reporting (C3), AR Collections (D1), and the Knowledge Base (E2) — sit in the integrated/financial-workpaper band; and the two flagships, the Field Examination Analyzer (A1) and the Comprehensive Cash Flow Model Builder (B1), sit at the upper end of that band (roughly **\$55,000–\$75,000**) given their reconciliation depth and audit-trail needs. Running several as one orchestrated pipeline is the multi-agent tier. Production agents that must be monitored, evaluated, and retuned after launch are covered by a **Managed AI — Agent Operations** retainer — **\$2,500–\$6,500/mo** depending on the number of agents in production, data-refresh cadence, and review SLA (distinct from the adoption-focused Managed AI tiers in Part 1).

Pool A — Field Exam & Due Diligence *(ARG Transaction Services)*

A1. Field Examination Analyzer & Collateral Roll-Forward *(recommended first build — flagship)*

Ingests borrower AR/AP agings, inventory listings, GL, and bank statements and produces the standard working-capital exam analysis: AR collateral roll-forward, turnover, gross and net dilution, credit-memo analysis, ineligibles, past dues, concentrations, aging trends, and collateral-to-cash reconciliation — plus inventory turnover, slow-moving and ineligible flags, and gross-margin testing. Outputs structured workpapers and a first-draft **Executive Report** in ARG's format.

- **Metric moved:** senior hours per exam; exam turnaround time; exam capacity per associate
- **Illustrative value:** materially reduced prep/analysis time per exam, freeing senior staff for findings and judgment (validate in assessment)
- **Data:** borrower AR/AP agings, inventory detail, GL, bank statements, prior ARG workpapers

A2. Borrowing Base Certificate (BBC) & Eligibility Agent

Validates borrower-submitted BBCs against underlying detail, recomputes eligible collateral and advance rates, and flags ineligibles, concentration limits, over-advances, and formula errors. Can run on a recurring portfolio-monitoring cadence, not just at exam time.

- **Metric moved:** over-advance risk caught early; reporting-review time
- **Illustrative value:** faster, more consistent BBC review; earlier detection of collateral shortfalls

- **Data:** BBCs, AR/inventory detail, loan agreement eligibility rules

A3. Portfolio Monitoring & Covenant Early-Warning Agent

Continuously watches borrower reporting for negative trends, covenant-breach risk, and deteriorating collateral — the “enhanced scope” trigger ARG describes — and alerts the relationship team while there's still time to act. Turns periodic review into continuous signal.

- **Metric moved:** time-to-detection of troubled credits; lender exposure
- **Illustrative value:** earlier intervention windows; more enhanced-scope and advisory engagements surfaced
- **Data:** periodic borrower financials, BBCs, exam history, covenant terms

Pool B — Cash Flow Modeling & Advisory *(ARG Partners)*

B1. Comprehensive Cash Flow Model Builder

Assembles the full financial-statement forecast ARG Partners is known for — forecasted balance sheet with calculated cash, projected P&L with margins and EBITDA, working-capital roll-forward (AR, inventory, AP, line of credit), and borrowing-base availability — drawing heavily on the field exam and prior reports that already supply ~80% of the inputs. Produces a model summary with assumptions and time-sensitive sources & uses.

- **Metric moved:** model build time; speed from engagement to advice
- **Illustrative value:** compresses comprehensive-model assembly; converts the existing “80% already exists” into hours saved
- **Data:** field exam output, borrower financials, appraisals, contract/debt schedules

B2. 13-Week Cash Flow & Variance Agent

Builds and maintains the 13-week model used in CRO/CFO and bankruptcy engagements: weekly cash in/out roll, matching working-capital roll-forward and borrowing-base calculation, and weekly budget-to-actual plus budget-to-date variance with draft commentary.

- **Metric moved:** weekly reporting effort during live engagements
- **Illustrative value:** turns a recurring weekly grind into a reviewed draft; supports lender-approved budget reporting
- **Data:** comprehensive model, weekly actuals, bank data

B3. Restructuring Scenario & Sources-and-Uses Agent

Models the cost/benefit of restructuring options — refinance, recapitalization, asset sale, right-sizing, or wind-down — projecting liquidity under each so principals can show stakeholders a defensible path.

- **Metric moved:** scenario analysis turnaround; quality of stakeholder communication
- **Illustrative value:** faster, side-by-side scenario comparisons for lenders and owners
- **Data:** comprehensive model, asset appraisals, debt terms

Pool C — Appraisal & Asset Recovery *(Recovery · Industrial · Asset Recovery)*

C1. Inventory Appraisal (NOLV) Accelerator

Structures borrower inventory listings, itemizes slow-moving and obsolete stock, runs sales and gross-margin trend analysis, studies the trial balance for sale expenses, and computes recovery scenarios — feeding ARG's Benchmarking Risk Algorithm with fresh, structured data for the 60-second at-a-glance risk views.

- **Metric moved:** appraisal prep time; monitoring cadence between reappraisals
- **Illustrative value:** faster NOLV analysis and more frequent benchmark refreshes
- **Data:** inventory listings (SKU-level sales/margin/mix), trial balance, market pricing

C2. M&E Appraisal Comparable-Research Agent

Researches current market and auction comparables for machinery and equipment by industry to support FLV, OLV, FMV, and in-use/in-place opinions — accelerating both desktop and inspected appraisals while keeping the certified appraiser's judgment central.

- **Metric moved:** comp-research time per appraisal
- **Illustrative value:** faster turnaround on M&E appraisals across ARG's many industries
- **Data:** equipment lists, auction/market comp sources, prior appraisals

C3. Wind-Down & Liquidation Reporting Agent

Drafts the recurring reporting that wind-downs generate: DIP/cash-collateral budgets, budget-to-plan tracking, loan-to-value tracking during liquidation, and bankruptcy schedules and operating reports — first drafts for the responsible-party team to review.

- **Metric moved:** reporting effort during wind-downs; reporting timeliness
- **Illustrative value:** less manual report assembly across long engagements
- **Data:** liquidation budgets, sale proceeds, asset registers, court reporting templates

Pool D — AR Management

D1. AR Collections Intelligence & Ledgered-Receivables Agent

Indexes a company's collection days against industry averages to quantify the cost of slow turnover, surfaces systemic credit/process issues, prioritizes accounts, and drafts improved credit policies and procedures. For ledgered-receivable engagements, maintains the weekly/monthly AR roll-forward and full collateral-to-cash reconciliation — even when the company's legacy accounting system is no longer active.

- **Metric moved:** collection-day diagnosis speed; ledgered-receivable reconciliation effort
- **Illustrative value:** faster diagnostics and cleaner collateral clarity for stakeholders
- **Data:** AR agings, payment history, credit policies, industry benchmarks

Pool E — Firmwide knowledge & market

E1. Commodity Report Research & Drafting Agent

Gathers monthly commodity pricing, computes trailing trends, and drafts the recurring **ARG Commodity Report** for client distribution — turning a standing monthly production task into a reviewed draft, and making it easy to add new commodities clients request.

- **Metric moved:** monthly production time; consistency and reach of the report
- **Illustrative value:** a recurring deliverable produced faster and more consistently every month
- **Data:** commodity price feeds, historical report archive

E2. Engagement Knowledge Base (RAG) + Proposal Drafter

A searchable, permissioned memory across ARG's past exams, models, and appraisals so professionals instantly reuse prior work — operationalizing the “prior reports supply 80% of what we need” advantage — and a drafting assistant that produces scoped field-exam and advisory proposals and engagement letters from a short brief.

- **Metric moved:** time spent re-finding prior work; proposal turnaround
- **Illustrative value:** faster scoping and onboarding; institutional knowledge retained as staff change
- **Data:** historical engagement archive, templates, scope library

E3. Executive / Engagement Briefing Agent

A governed daily or weekly digest for ARG principals across all active engagements — borrowers trending negative, upcoming deadlines, covenant and collateral risks, and the exam/appraisal pipeline — with every figure traceable to its source.

- **Metric moved:** principal time; decision latency across the portfolio
- **Illustrative value:** hours/week returned per principal; fewer items falling through the cracks
- **Data:** the agents above, engagement systems, calendar/deadlines

---

Flagship deep-dive: the Field Examination Analyzer (A1)

ITECS recommends the field-exam analyzer as ARG's first build because it sits at the **front door of almost everything else** — a clean, structured exam feeds the cash flow model (B1), the BBC agent (A2), and portfolio monitoring (A3) — and it targets ARG's scarcest resource: senior examiner time.

The shift it delivers: today a senior associate spends hours normalizing a borrower's AR/AP agings, inventory, GL, and bank statements before the analytical judgment begins. The analyzer does that assembly — roll-forwards, dilution, turnover, ineligibles, concentrations, reconciliations — and hands the examiner a structured workpaper set and a first-draft Executive Report. The professional moves straight to *interpretation and findings*, which is what lenders actually pay ARG for.

Why it's the right first build:

- **Highest volume, most standardized** — the standard-scope working-capital exam repeats across every engagement, so one well-built agent pays back across the whole portfolio.
- **Data already arrives** — borrowers and prior ARG workpapers supply the inputs; no greenfield integration required to start.
- **Feeds the franchise** — the same structured output shortens the comprehensive cash flow model, which ARG notes is ~80% built once a fresh exam exists.
- **Governed by default** — the agent drafts; the examiner reviews, adjusts, and signs. Every figure traces back to its source document.

What leadership sees: faster exam turnaround for lenders, more exams handled by the same senior bench, and consistent, audit-ready workpapers — without taking the senior professional out of the judgment seat.

Indicative build & run: an Agent Discovery & Technical Spec ($4,500–$7,500), then the analyzer as an integrated financial-workpaper agent (~$55,000–$75,000), then a Managed AI — Agent Operations retainer ($2,500–$6,500/mo) to keep it tuned — phased so ARG sees a working prototype before committing to the full production build.

How ITECS calculates and stands behind ROI

- **Illustrative now, validated in the assessment.** Figures here are framed as time and throughput leverage, not promised dollars. The readiness assessment replaces them with ARG's real engagement economics and a defensible baseline.
- **The right ROI metric for ARG is senior capacity.** ARG bills senior judgment; value shows up as hours freed, faster turnaround for lenders, and more engagements per professional — plus quality and consistency of workpapers.
- **Phased and gated.** Orientation and the readiness assessment come first; a first agent build is scoped only after the economics are ranked, so spend follows proven payback.
- **Confidentiality is the design constraint, not an afterthought.** Everything is built for a firm that handles lender and distressed-borrower data under strict confidentiality.

Governance & security — built for confidential financial work

- **Read broadly, act carefully.** No autonomous client communications, no changes to collateral records, no finalized valuations or reports without a professional's explicit review and approval.
- **Role-based access** aligned to ARG's identity provider; agents only see what their human owner is permitted to see.
- **Audit-ready by default.** Every recommendation preserves its assumptions, source data, and the reviewer's decision — essential when work product may be examined by lenders, courts, or in litigation support.
- **Confidential boundary.** Business/enterprise AI tiers that contractually isolate ARG's data and never train on it; credentials and integration secrets managed outside prompts.
- **Credibility:** delivered on ITECS's managed-security foundation (SOC 2 Type II, ISO 27001, and CMMC-aligned practices).

The path from here

4. **Executive AI Orientation** ($3,500) — a half-day session with ARG principals to align on where AI helps and where the guardrails go.
5. **AI Readiness Assessment** ($6,500) — map ARG's workflows and data, scan for shadow AI, and rank the agent menu by payback.
6. **Foundation** — AI Acceptable Use Policy ($4,500) and Shadow AI Discovery ($5,500) to make adoption safe firmwide.
7. **First build** — start with an Agent Discovery & Technical Spec ($4,500–$7,500), then build the Field Examination Analyzer as an integrated financial-workpaper agent (~$55,000–$75,000), and expand into the cash flow model builder and portfolio monitoring.
8. **Sustain** — Internal AI Champion program, a Managed AI Services retainer for adoption, and a Managed AI — Agent Operations retainer to keep production agents monitored and tuned.

ARG and ITECS already have a working relationship and a security posture in place — which means a shorter, lower-risk path to value than a cold start with a new vendor.